

## Relational Databases

- 1) lotta DBs out there
- 2) 2 major categories dependin on struc imposed on these DBs

### ① Relational a.k.a SQL DB (eg:- POSTGRESQL)

(i) A type of DB that imposes on the data stored in it, a tabular like structure (data stored as table)

(ii) Rows → instances of the entities that the resp. tables represent

Columns → Attributes of the entities

eg:- table → payments record for a shop  
 row → each customer payment details  
 such as date, amount, transaction medium  
 column → the attr. ↑ ↑ ↑

Cust. name	Date	Amt	Transac <sup>n</sup> medium
A	17/9	10	Credit Card
B	15/9	45	Cash
C	14/9	3	Melbankin
D	13/9	87	Cash

(iii) Schema → Define the shape of the table i.e. it will basically specify what attr (col) each entity (row) will have

eg:- Schema for a book library

Each row → book name  
 For each row, cols → borrowed on, returned on, details of payment, condition of book

② Non relational DB aka NoSQL (eg! - mongo DB, Google Cloud Datastore)

(i) DBs don't impose tabular struc. on data stored in them. Sometimes they may impose some kinda other struc. but def by not table

(ii) Flexible, less rigorous

SQL → Structured Query Lang.

i) Most relational DB (rDB) support SQL

ii) SQL is used to perform <sup>powerful</sup> complex queries in rDBs

iii) Because rDBs support SQL, they ~~support~~ come with powerful querying capabilities and this is also the reason why most people choose rDBs over non-rDBs for part of their sys.

iv) Why SQL > Python, JS

Cuz for Python, JS you gotta load data in mem. to perform these kinda queries

So when we talk abt large scale distri. sys, we got Tbs of data and we can't do this trivially



## SQL DBs aka RDBs

- 1] Must use ACID transac<sup>n</sup>s → transac<sup>n</sup> supporting ACID props.
- Atomicity   Consistency   Isola<sup>n</sup>   Durability

i) Atomicity - If a transac<sup>n</sup> consists of multiple sub-op then these sub-op are gonna be considered as 1 unit so they'll either all succeed or all fail

eg: - \$ funds transfer in banks

→ Sub-opn 1 → Reduce money in person 1's acc.  
→ Sub-opn 2 → Increase money in person 2's acc.

ii) Consistency  
(aka Strong Consistency)

↓  
dB is never invalid and never stale

① Any transac<sup>n</sup> in dB is gonna abide by all the rules in the dB → dB is never invalid  
② Any future transac<sup>n</sup> in dB, is gonna consider any past transac<sup>n</sup>s in dB  
i.e. no 'stale' state in dB where 1 transac<sup>n</sup> has exec.<sup>d</sup> but another transac<sup>n</sup> doesn't know it has exec.<sup>d</sup>  
→ dB is never stale

iii) Isola<sup>n</sup> - Multip. transac<sup>n</sup>s can occur at the same time but, under the hood, they will actually hr been exec.<sup>d</sup> as if they had been done sequentially (i.e. put in a queue) one-by-one

iv) Durability - When you make a transac<sup>n</sup> in a dB, the effects of that transac<sup>n</sup> in the dB are permanent

Date: \_\_\_\_\_  
M T W T F S S

ie. data stored in the db is effectively stored in disc (not mem.)

## Database Index (DBI)

A DBI is a data struct that improves the speed of data retrieval ops on a db at the cost of additional writes and storage space to maintain the index data struct

Consider this eg for bank passbook

Month	Date	Amount credited	Amt debited
January	28 <sup>th</sup>	10000	—
Feb	7 <sup>th</sup>	5000	—
March	15 <sup>th</sup>	700	—
April	27 <sup>th</sup>	2000	—

Say we wanna figure out on which day of the yr was the highest amt. credited. (Ans: Jan 28<sup>th</sup>)

(i) Conventional Meth. :- Traverse thru entire book  
thru entire db to find that amt  
 $O(N)$

(ii) Using DBI :- Auxiliary data struct. created  
ie. optimized for fast searching  
on a specific attr. (col.) in  
the table.

Say our DBI = table that has all 'credited amts'  
stored in sorted order ( $O(N \log N)$ ) and



each of these amts pts. to the relevant record (row) in the main dB table  
thus, we get our largest amt by simply going to the end of the dBI table

### simple way of lookin at dBI

- auxiliary data struc of main dB that speeds up getting data from main dB (read ops. faster)
- auxiliary data struc so ~~so~~ extra space reqd. for storage
- whenever you write data to main dB, ~~data~~ you also gotta write to dBI (write ops. slower)
- in practice, we create an idx on 1 or multip cols. in our main dB

### NOTE : Eventual Consistency

- A consistency model which is unlike Strong Consistency
- In this model, reads might return a stale view of the sys.
- An eventually consistent ~~db~~ datastore will give guarantees that the state of the data will eventually reflect writes within a time period (eg: - 10 mins, 30 days)
- eg: - Google Cloud Datastore